Gavin Valencia
Assignment 7
12/14/18

Question 1

1. i) cn33
   ii) Tesla P100-PCIE

2. 6.0

3. nvcc -arch=compute_60 -code=sm_60 -lcuda main.cu -o main

Question 2

I was unable to get Monsoon working, I keep getting my permission denied. And, since I started too late I did not have time to ask for help. However, I did write the code to time it, and it would be compiled as such:
srun nvcc -arch=compute_60 -code=sm_60 -lcuda -Xcompiler -fopenmp question2_gdv4.cu -o question2_gdv4

In addition, I'd imagine with a high enough N, the computation in the GPU will take a relatively small amount of time, but the memory transfer will take less time than the overall computation on the CPU, making it the more efficient of the two.

Question 3

1)
I was unable to get these running either. However, I bet as N increases the GPU becomes much more efficient.

2)
My algorithm works by trimming off any duplicate points and then sending them to the GPU. However I was unable to run it.